

Himanshu Laddhad

765-414-2728 | himanshuladdhad11@gmail.com | LinkedIn | GitHub | Portfolio | Austin, TX

SUMMARY

Data Engineer with 3 years of data experience building tested pipelines on SAP Datasphere, dbt, and Python - cutting manual reporting 20+ hours/week and stopping bad data before it reaches KPIs.

Certifications: Tableau Desktop Specialist, Microsoft AI-900, AWS CLF-C02, DataCamp: Data Scientist Associate & SQL Associate

Languages & Libraries: SQL, Python, PySpark, PostgreSQL, DuckDB, ETL Pipelines, ELT Pipelines, Pandas, NumPy, Scikit-learn, XGBoost, PyTorch
AI, MLOps & Platforms: dbt, Apache Airflow, Snowflake, Databricks, BigQuery, SAP Datasphere, Kafka, Amazon Web Services, Microsoft Azure, Git, Power BI, Tableau, Delta Database, Relational Database, NoSQL, MLflow, Power Automate, LangGraph, Large Language Models, Retrieval-Augmented Generation, QLoRA, PEFT, LangChain, AWS Sagemaker, GraphDB, Microsoft Excel, Microsoft PowerPoint.

Methods & Domain: Data Pipeline Architecture, Data Warehousing, Dimensional Modeling, Schema Design, Incremental Loading, Orchestration, Data Quality Testing, Data Contracts, Machine Learning, A/B Testing, Agile

EDUCATION

Purdue University, Daniels School of Business

West Lafayette, IN

Master of Science in Business Analytics and Information Management, GPA: 3.7/4.0

August 2025 – July 2026

Sardar Vallabhbhai National Institute of Technology

Surat, India

Bachelor of Technology, Mechanical Engineering, GPA: 3.6/4.0

August 2019 – May 2023

PROFESSIONAL EXPERIENCE

PartnerLinQ

West Lafayette, IN

AI Engineer - Consulting

January 2026 – May 2026

- Architected a **RAG**-grounded data transformation pipeline routing messages through Groq to a GPT-4.1 tool-calling model, retrieving from a 400+ file XSLT corpus - cutting mapping build time from **3-5 weeks to 3-4 hours** and letting new partners onboard in hours, not weeks.
- Engineered a **QLoRA** fine-tuning pipeline training an SLM (Qwen) on transform generation, modification, repair, and explanation, grounded by the same retrieval layer at **68% first-pass Altova validity** vs. the GPT-4.1 pilot's 78%, keeping partner mapping data in-house.
- Delivered the **Streamlit evaluation harness** fronting both pipelines, executing generated XSLT against user-supplied test cases in-session, letting engineering team validate and trust model output and integrate it into their platform as a chat agent or assistant.

JCB India Ltd

Pune, India

Data Engineer - Business Intelligence and Quality

February 2024 – July 2025

- Consolidated four event streams from a delta database into one **SAP Datasphere** pipeline (inbound, harmonized, curated) on a single serial join key at **35K rows/day** - replacing manual consolidation for 10+ quality, engineering, and operations consumers.
- Hardened the pipeline with **dbt tests** and custom SQL drift checks that quarantine failed rows with typed exception codes - stopping bad data before it could distort KPIs or supply ordering.
- Scheduled hourly incremental refresh from SAP HANA into **Power BI** for 15+ KPIs across 3 stakeholder groups - cutting **20+ hours/week** of manual reporting, redirecting analyst time to data quality.
- Orchestrated severity-tiered exception alerting through **Power Automate** to Teams adaptive cards with writeback acknowledgement, cutting **mean time-to-acknowledge 15%** and easing the backlog behind delayed production commitments.

JCB India Ltd

Pune, India

Data Engineer - Supply Chain

May 2023 – February 2024

- Instituted **freshness and sanity checks** as the first pipeline steps after a stale input produced silently wrong forecasts - halting bad data before it reached procurement and averting an estimated **\$5K-\$20K** in wrong parts orders per incident.
- Led a 3-person rebuild of the demand-forecasting pipeline, migrating from Excel to **XGBoost** in Python on SAP BW, engineering lag and volatility features across 500+ SKUs and 4 product lines - cutting **forecast error 18%** and reducing excess/deficit inventory.
- Built a Python **BFS traversal** over a multi-level component hierarchy, exploding SKU forecasts into per-component monthly demand for procurement - resolving shared-component conflicts and holding **98% on-time delivery** vs the 95% target.
- Deployed Python time-series monitoring flagging demand deviations **4-6 weeks ahead** across 6+ markets, letting planning teams correct course via ocean freight instead of air - roughly **15x cheaper**.

Bosch

Nashik, India

Data Analyst (Operational Analytics)

April 2022 – August 2022

- Structured a Pareto data analysis over 9 months of operational time-series data in Python, ranking **140+ exception categories** by contribution to downtime - lifting **process efficiency 10%**.
- Integrated corrective-action tracking against the ranked categories across procurement, quality assurance, and operations - cutting **exception frequency ~90%** and **reprocessing rate 14%**.

PROJECTS

- Causal Promotion System** - A campaign attribution tool separating true promotion lift from coincidence, segmenting customers by discount responsiveness using **CATE meta-learners** (Python, EconML, DuckDB) on A/B test data
- Customer Churn & Lifetime Value** - A **DuckDB SQL feature pipeline** over **410M log rows** for 6.7M users (two-pass aggregation to fit memory, leakage checks) feeding XGBoost/LightGBM churn and lifetime-value models, projecting \$1.26M in expected value.
- Bankruptcy Prediction Ensemble** - A financial risk classifier flagging bankruptcy from 64 financial ratios, combining XGBoost, LightGBM, and CatBoost (Python, Optuna) into an ensemble reaching **0.90 validation ROC-AUC**.

AWARDS AND AFFILIATIONS

- JCB - Excellence in Planning and Efficiency (2025), Rising New Joinee Award (2024)
- Indian Institute of Management Ranchi - Research Intern, Market & Competitive Analysis of EV Landscape
- Purdue University Graduate Teaching Assistant, AI for Business (June-July 2026)